

# Digital Forensics

## Lecture Week 8

# Linux artifacts

Nelson Chapters 3 and 7  
Readings

# Linux forensics

- I need to run forensics Linux tools
- How?
- “Our Linux Server has been infected”
- Why and How?

# Linux boxes

- Where do I get a Linux box to examine?
- Use WSL on Windows 10
  - Note: A VM will not show kernel activity as this is on the host
- Use a terminal window on MacOS
- Use a Linux VM
- Hire a Linux Server on AWS for \$6 a month

# A Linux Server Case Study

- A company used an ftp server to distribute software updates
- One day it stopped working as the OS had been erased
- The IDS logs had three high priority events
- These indicated a WU\_FTP attack and a file upload to the ftp server
- Log file analysis showed the file to be mount.tar.gz
- Using **strings** and **grep** on output of log files showed well known root kit files being installed
- A backdoor was installed, but the Firewall prevented files been sent out of the network
- <http://www.linuxsecurity.com/content/view/117644/>

# Objectives

- To understand basic Linux
- To understand the Linux system
- To locate volatile evidence
- To locate non volatile evidence
- To examine log files

# Linux OS components

- The **kernel**
  - talks to the CPU and Hardware
  - modular
- The **operating system tools** (GNU)
  - compilers and libraries
  - the shell and command line tools
- The **user interface (environment)**
  - gui
  - Touch
- The **applications**
  - Do useful things like run a web server

# Linux OS Distributions

- Assemble the components in a particular flavour
- **Debian GNU**(1993) Descendants
  - Knoppix
  - Ubuntu (2004)
    - Backtrack, Kali
- **Red Hat** (1994-2004) Descendants
  - Fedora (free)
  - CentOS (2003)
  - Red Hat Enterprise (RHEL) (mainly in USA)
- **Open SUSE** (Novell) 1994
  - Mainly in Europe

# Debian Releases – Aug 2020

4.0	Etch	2007-04-08	Etch, the Etch-A-Sketch
5.0	Lenny	2009-02-14	Lenny, the binoculars
6.0	Squeeze	2011-02-06	Squeeze toy aliens
7	Wheezy	2013-05-04	Wheezy the penguin
8	Jessie	2015-04-26	Jessie the cowgirl
9	Stretch	2017-06-17	Rubber octopus from Toy Story 3
10	Buster	2019-07-06	Andy's pet dog
11	Bullseye	Not yet released	Woody's horse
	Sid	"unstable"	The next door neighbour



# OS Distribution Packages

- Download as a DVD image (.iso) file or a VM
- **Debian**
  - Includes Kali, Knoppix and Ubuntu
  - use the **apt** package manager
- **Fedora**
  - includes CentOS, RHEL
  - use the **yum** package manager
- Distro ranking and latest news
  - <http://distrowatch.com/>

Page Hit Ranking	
Last 6 months	
Rank	Distribution
1	<a href="#">MX Linux</a>
2	<a href="#">Manjaro</a>
3	<a href="#">Mint</a>
4	<a href="#">Ubuntu</a>
5	<a href="#">Debian</a>
6	<a href="#">elementary</a>

# Linux user interfaces

- The gui interfaces
  - Use X Windows (also called X11) for bit mapped displays
  - KDE (Windows like)
  - Gnome (Apple like)
  - Blackbox (minimal X11)
- The cli Interface
  - command line interface
  - like Cisco cli and windows cmd
  - use a shell, usually `bash`

# Linux applications

- Usually issued as a **package**
- Downloaded and installed using a package manager
- The package manager also specifies a package format
- **RPM** for Red Hat
- **APT** for Debian using the **dpkg** format
- **Licencing**
  - all code must be compatible with version 2 of the GNU General Public License (GPLv2)
  - all code must be signed
  - All commands support the **-help** option

# Linux and Malware

- Small companies often use Linux as the server platform
  - Web server, ftp server, mail server, db server
  - Thus a target for web application based attacks
- Attackers often use Linux as their C & C Server
  - Many Linux servers are not well protected
  - Some infections last for years undetected


# Linux Security Features

- iptables – the Linux Firewall
- <http://en.wikipedia.org/wiki/Iptables>
- Virus Scanners
  - not common on Linux, [clamav](http://www.debianhelp.co.uk/clamav.htm) is one
  - <http://www.debianhelp.co.uk/clamav.htm>
- Use of private key Certificates for networking
  - openssl, ssh, openvpn, etc
- Sudo (superuser do)
  - privilege escalation as required
- <https://wiki.debian.org/sudo>

# Objectives

- To understand basic Linux
- To understand the Linux system
- To locate volatile evidence
- To locate non volatile evidence
- To examine log files


# The Linux Kernel



# The Linux Kernel Archives



Protocol	Location
HTTP	<a href="https://www.kernel.org/pub/">https://www.kernel.org/pub/</a>
GIT	<a href="https://git.kernel.org/">https://git.kernel.org/</a>
RSYNC	<a href="rsync://rsync.kernel.org/pub/">rsync://rsync.kernel.org/pub/</a>



mainline:	<b>5.8</b>	2020-08-02
stable:	<b>5.8.1</b>	2020-08-11
stable:	<b>5.7.15</b>	2020-08-11
longterm:	<b>5.4.58</b>	2020-08-11
longterm:	<b>4.19.139</b>	2020-08-11
longterm:	<b>4.14.193</b>	2020-08-07
longterm:	<b>4.9.232</b>	2020-07-31
longterm:	<b>4.4.232</b>	2020-07-31
linux-next:	<b>next-20200814</b>	2020-08-14



# The WSL Kernel

- Here the Kernel is modified to run on Windows
  - Windows Subsystem for Linux (WSL)
- Looking at the kernel inside Linux will fail

```
$ whatis uname
uname (1)      - print system information
uname (2)      - get name and information about current kernel

$ uname -r
4.4.0-20190-Microsoft      # Returns Windows version of Linux
```

- You need to ask **wsl**

```
C:\Users\graha>wsl --update --status
Windows Subsystem for Linux was last updated on 24-Jun-20
WSL automatic updates are on.
Kernel version: 4.19.121.1
```

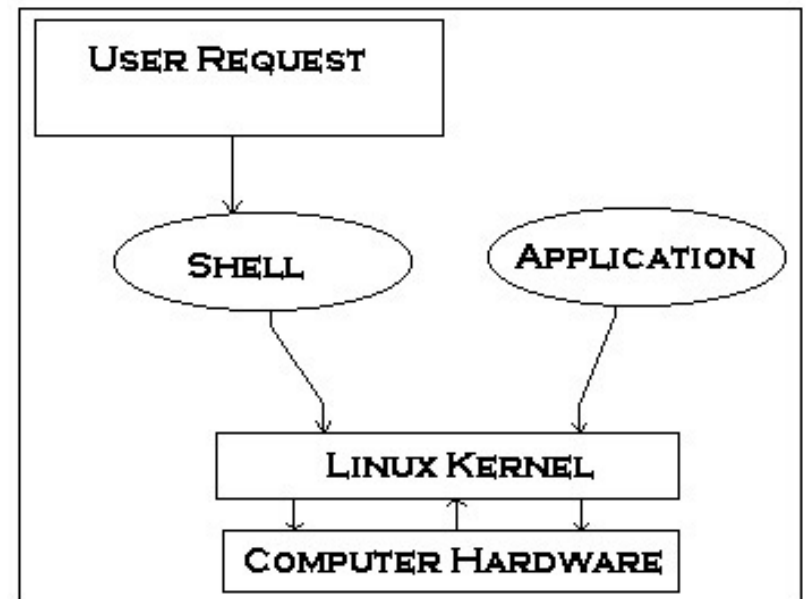
- More about **wsl** in readings

# The Linux system

- A bootloader
  - GNU GRUB or LILO
  - loads the kernel from a file into RAM
- init
  - the top of the process tree, launched by the kernel
  - launches other processes
- libraries
  - contain code used by other processes
  - GNU C library ([glibc](#))
  - like dlls in Windows

# The Linux system #2

- The GNU C Compiler gcc
  - compiles and links the programs used by Linux
- the user interface
  - a shell or gui
- Shells
  - the CLI to talk to the kernel
  - takes commands from stdin
  - supports regular expressions
  - keeps a history



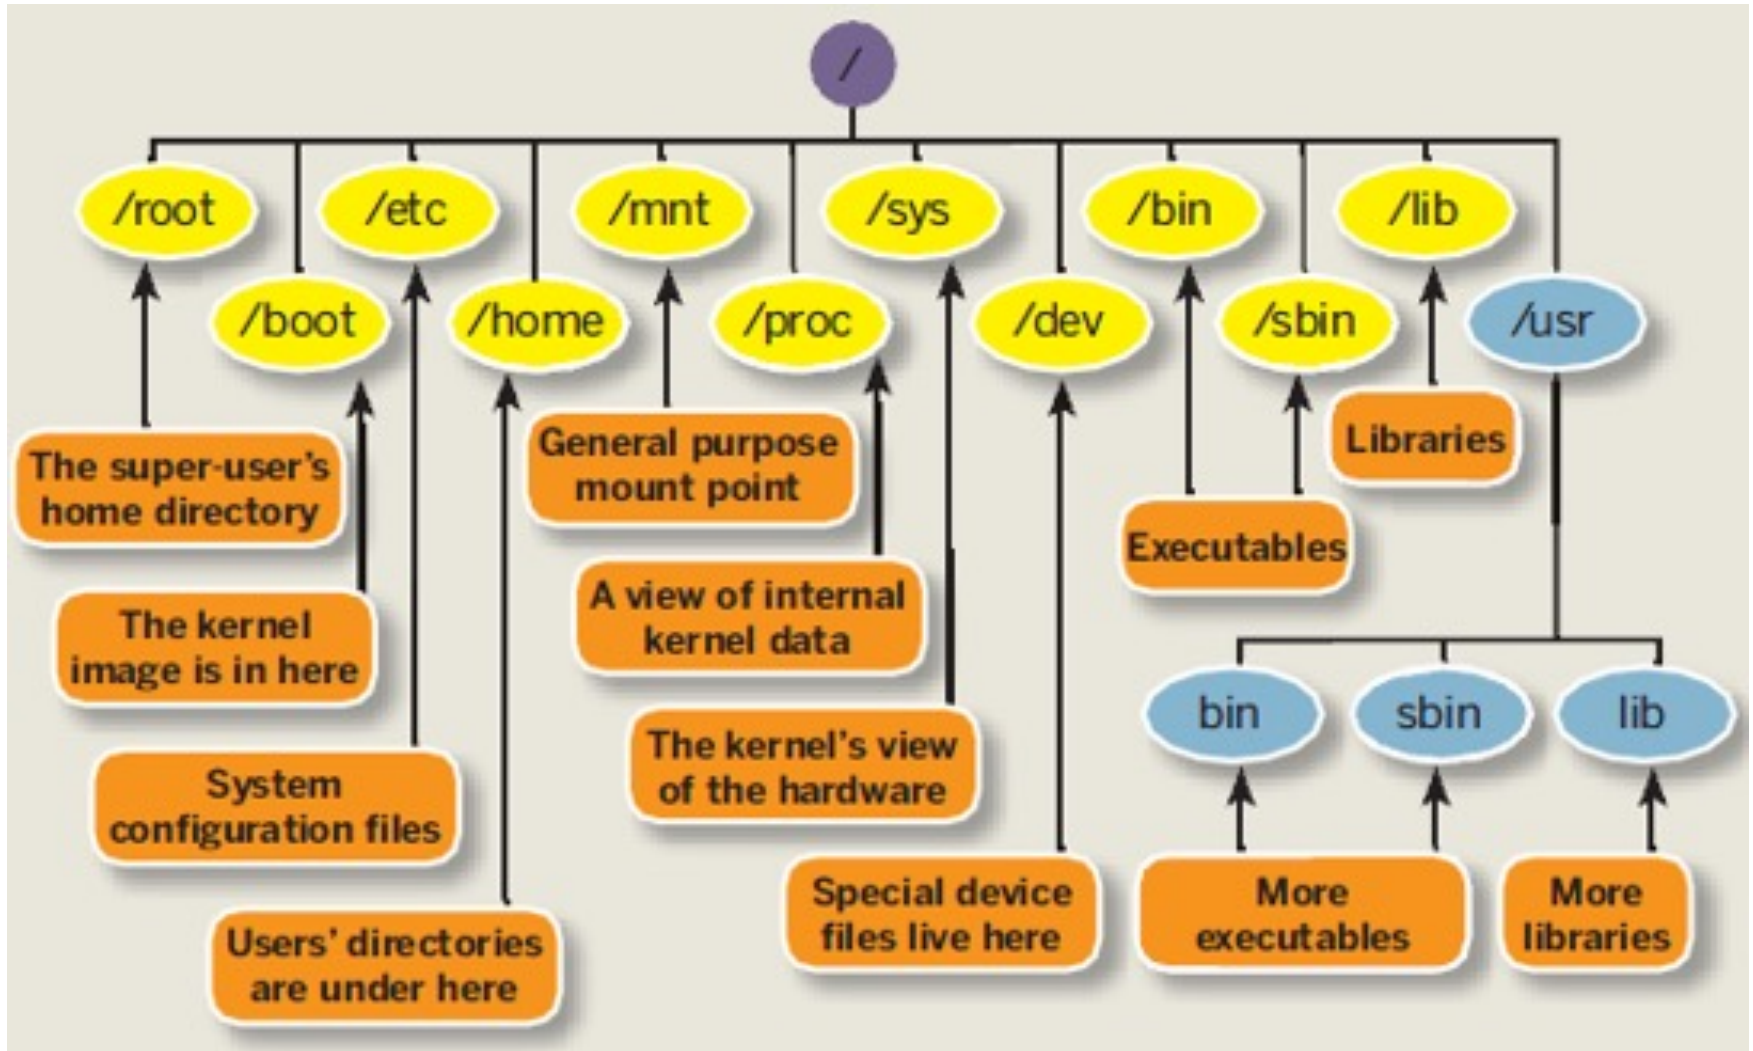
# Shell versions

- **bash**
  - Bourne Again Shell, part of GNU Linux
- **cs**
  - The C Shell, uses C syntax
- **sh**
  - the system default shell
  - often the **Dash** shell for speed (Derived Almquist Shell)
- Other shells
  - Korn Shell, Tenex C Shell

# Linux File structure

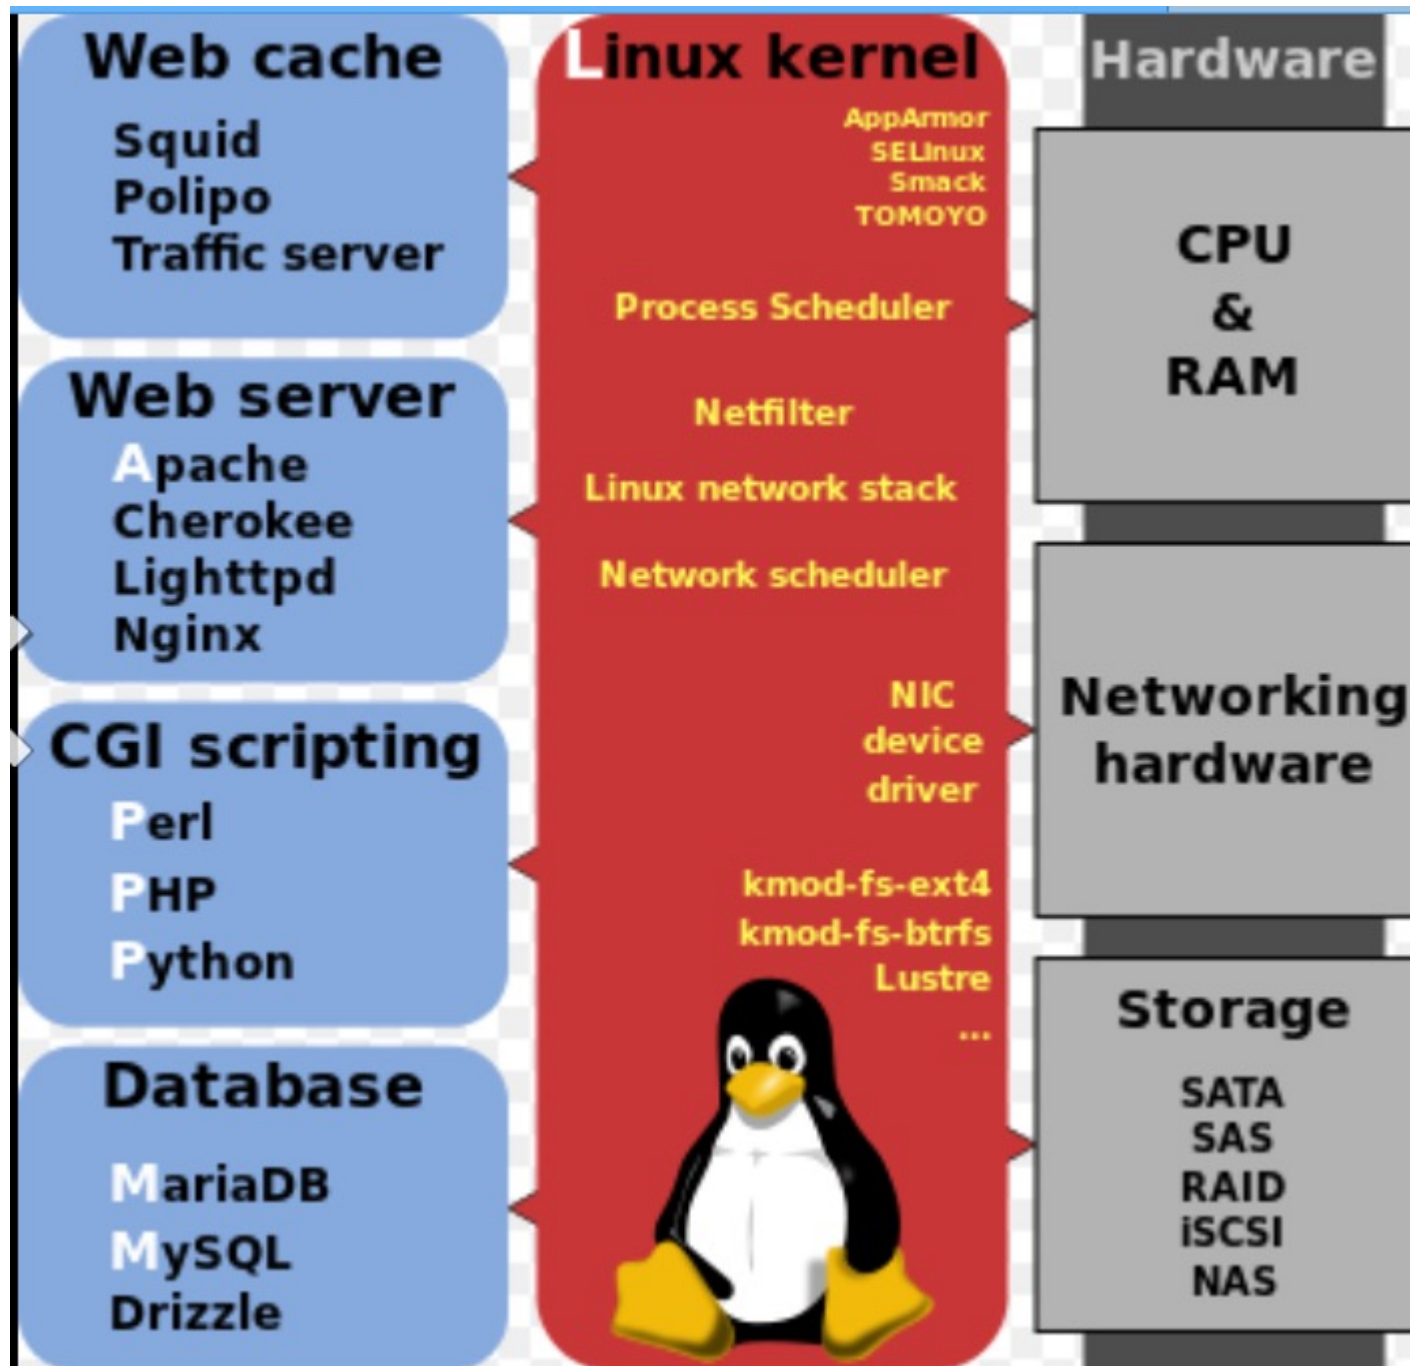
- / - root
- /bin – binaries (executables)
- /lib – libraries (dlls)
- /dev – devices (disks, ram, usb)
- /etc – config files (registry)
- /home – user folders
- /var – logs, swap files, mailboxes, caches

# Linux File structure #2



# LAMP

- Many small Linux Servers run LAMP
- An open source solution stack that works
- High Availability, Heavy Duty
- Linux (L) OS
  - Debian and Ubuntu share 60% of Linux web servers
- Apache2 (A) Web Server (56%) - some nginx (25%)
- MySQL (M) Database – also PostgreSQL
- PHP (P) Dynamic Web Pages – also Perl, Python





# Objectives

- To understand basic Linux
- To understand the Linux system
- **To locate volatile evidence**
- To locate non volatile evidence
- To examine log files

# Linux forensic shell commands

- Many tools are the same as for Windows
  - netstat, ifconfig (ipconfig), date, ping
- Some are Unix based
  - ps (process), df (mount points), du (disk usage)
  - uname (OS version), w (logged on users)
- Some require root privilege (sudo)
  - fdisk, crontab, viewing logs
- Some will not work on VMs such as WSL
- See the labs for details

# sudo

- SuperUser Do (sudo)
- Ubuntu supports restricting dangerous commands to the SuperUser called **root**
- To run a root command as User just prefix **sudo**
- `sudo ifconfig`

```
group11@kali:~$ sudo ifconfig  
  
We trust you have received the usual lecture from the local System  
Administrator. It usually boils down to these three things:  
  
#1) Respect the privacy of others.  
#2) Think before you type.  
#3) With great power comes great responsibility.  
  
[sudo] password for group11:
```

# who can sudo?

- root is **su** so root can sudo
- To see who else can **sudo** look at the sudo group

```
root@kali:~# cat /etc/group | grep sudo  
sudo:x:27:group11
```

# Volatile evidence

- who is logged on remotely?

- `w` nice and simple

- running processes

- `ps -af` local processes

- `ps -Af` system processes

- services

- `service -s status-all`

- `ls /etc/init.d` startup processes

```
$ service --status-all | grep +  
[?] apport  
[?] cryptdisks  
[?] hwclock.sh  
[+] irqbalance  
[+] iscsid  
[+] lvm2-lvmetad  
[+] lvm2-lvmpolld  
[?] networking
```

```
$ ls /etc/init.d  
acpid  
apparmor  
apport  
atd  
bootmisc.sh  
checkfs.sh
```

# Folders of forensic interest

- /etc/passwd – user names
- /etc/shadow – password hashes
- /etc/init.d – services
- /var/www – web server pages
- /var/log – log files
- /var/lib/mysql – database data files also PostgreSQL

# Objectives

- To understand basic Linux
- To understand the Linux system
- To locate volatile evidence
- **To locate non volatile system evidence**
- To examine non volatile log files

# non volatile Evidence

- Linux keeps memory information in [/proc](#)
- This is a virtual folder – a link to memory
- [/proc/cmdline](#) show how the boot image is loaded
  - the boot file name
- [/proc/cpuinfo](#) shows some CPU details
  - CPU model, speed and flags
- [/proc/meminfo](#) shows the memory manager details
  - total memory, swap file details, Virtual Memory details



# /proc/ examples

```
$ cat /proc/cpuinfo | grep processor
```

```
processor      : 0
```

```
processor      : 1
```

```
processor      : 2
```

```
processor      : 3
```

```
$ cat /proc/cpuinfo | grep 'model name' | uniq
```

```
model name     : Intel(R) Core(TM) i5-7400 CPU @ 3.00GHz
```

```
$ cat /proc/meminfo | grep Mem
```

```
MemTotal:      8259836 kB
```

```
MemFree:       4543520 kB
```

# Objectives

- To understand basic Linux
- To understand the Linux system
- To locate volatile evidence
- To locate non volatile system evidence
- To examine non volatile log files

# Linux log files

- The main logs are in **/var/log** (some not used in a VM)
  - (Fedora has different logs to Debian)
- **apache2**                      web server access
  - access.log                      – web visitors
  - ssl-access.log                  – ssl visitors
- **auth.log**                      authentication (password)
- **mail.log**                      mail
- **mysql.log**                    database
- **lpr.log**                      line printer
- **syslog**                      app logging

# journalctl

- Modern Linux uses **systemctl** instead of init.d to launch processes. (may not be visible in a VM)
- systemctl has its own log – **journalctl**
- Check with **hostnamectl** for device details

```
ubuntu$hostnamectl
  Static hostname: ip-172-26-5-152
        Icon name: computer-vm
        Chassis: vm
        Machine ID: 0eaed5dbe31548e38562f54383ed1376
        Boot ID: 8fbc0ef2f0fa497fb36226922559bc0a
  Virtualization: xen
  Operating System: Ubuntu 16.04.3 LTS
        Kernel: Linux 4.4.0-1049-aws
  Architecture: x86-64
```

# journalctl options

Reference – not assessed

- `journalctl | grep -i GHz`      # look for CPU info
- `journalctl -u ssh`      # look for unit
- `journalctl -t sshd`      # look for syslog identifier
- `journalctl -t dhclient | grep bound`      # to interface
- `journalctl -p 3`      # set warning level
- `journalctl --since -1w`      # open archive, w = week

# Useful events #1 - User Logon

- Forensics Sequence
- a) Identify the usernames - SID
- b) Find the user sessions – time and PID
- c) Find the log entries

## a) Identify the usernames

- Users are registered in the `/etc/passwd` file
  - many are system users with no shell
  - people have a shell called `bash`
  - Bourne-again shell (bash) replaces the original Bourne shell

```
root@kali:~# whatis cat
cat (1)          - concatenate files and print on the standard output
```

```
root@kali:~# cat /etc/passwd | grep bash
root:x:0:0:root:/root:/bin/bash
postgres:x:114:125:PostgreSQL administrator,,,:/
group11:x:1000:1000:,,,:/home/group11:/bin/bash
```

## b) Find the user sessions

- To see the logged in sessions use **last**

```
root@kali:~# whatis last
last (1)          - show listing of last logged in users
```

```
root@kali:~# last root
root    pts/2    192.168.198.1    Tue Jan  9 23:18    still logged in
root    pts/1    192.168.198.1    Tue Jan  9 23:00    still logged in
root    pts/0    :0.0            Wed Sep 21 00:54 - down    (00:01)
root    tty7     :0              Wed Sep 21 00:54 - down    (00:01)
```

```
root@kali:~# last group11
group11 pts/0    :0.0            Tue Jan  9 22:40    still logged in
group11 tty7     :0              Tue Jan  9 22:40    still logged in
```



# System reboots

- A special user called **reboot**

```
root@kali:~# last reboot
reboot    system boot  3.7-trunk-686-pa Tue Jan  9 22:40 - 23:59 (01:19)
reboot    system boot  3.7-trunk-686-pa Wed Sep 21 00:54 - 00:55 (00:01)
reboot    system boot  3.7-trunk-686-pa Sat Sep 17 04:09 - 04:10 (00:00)
reboot    system boot  3.7-trunk-686-pa Sat Sep 17 04:05 - 04:09 (00:03)
```

- See **system** with **-x**

```
$last -x -n 20
...<output snipped>
runlevel (to lvl 5)  4.4.0-1049-aws  Sun Feb 11 08:25  still running
reboot    system boot  4.4.0-1049-aws  Sun Feb 11 08:25  still running

$last --help | grep '\-n,'
-n, --limit <number> how many lines to show
$last --help | grep '\-x,'
-x, --system          display system shutdown entries and run level changes
```

## c) Find the log entries

- `cat /var/log/auth.log | grep gdm`
  - Gnome display manager (gdm) is a **local** login

```
root@kali:~# cat /var/log/auth.log | grep gdm3:session | grep group11
Jan  9 22:40:29 kali gdm3][3129]: pam_unix(gdm3:session): session opened for user group11
```

- To see **remote** logins, use ssh

```
root@kali:~# cat /var/log/auth.log | grep sshd

Jan  9 22:40:07 kali sshd[3315]: Server listening on 0.0.0.0 port 22.
Jan  9 22:40:07 kali sshd[3315]: Server listening on :: port 22.
Jan  9 23:00:43 kali sshd[3773]: Accepted password for root from 192.168.198.1 port 2047 ssh2
Jan  9 23:00:43 kali sshd[3773]: pam_unix(sshd:session): session opened for user root by (uid=0)
Jan  9 23:18:15 kali sshd[5430]: Accepted password for root from 192.168.198.1 port 2077 ssh2
Jan  9 23:18:15 kali sshd[5430]: pam_unix(sshd:session): session opened for user root by (uid=0)
```

## Useful events #2 - dhcp activity

- a) Get the current ip address details – MAC address
- b) Find the log entries – MAC or ip

## a) Get the current ip address details

- Use **ifconfig**

```
root@kali:~# ifconfig
eth0      Link encap:Ethernet  HWaddr 00:0c:29:b3:49:7b
          inet addr:192.168.198.128  Bcast:192.168.198.255  Mask:255.255.255.0
          inet6 addr: fe80::20c:29ff:feb3:497b/64 Scope:Link
          UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
          RX packets:4942 errors:0 dropped:0 overruns:0 frame:0
          TX packets:3747 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:1000
          RX bytes:466146 (455.2 KiB)  TX bytes:2670120 (2.5 MiB)
          Interrupt:19 Base address:0x2000
```

## b) Find the log entries

date and time, ip address and dhcp server

- We look for an ip address **offer** from the dhcp server
- `cat /var/log/syslog | grep -i dhcponffer`

```
root@kali:~# cat /var/log/syslog | grep -i dhcponffer
Sep 17 04:05:21 kali dhclient: DHCPOFFER from 192.168.28.254
Sep 21 00:54:08 kali dhclient: DHCPOFFER from 192.168.3.254
Jan  9 22:40:05 kali dhclient: DHCPOFFER from 192.168.198.254
```

- Then we look for a following **preinit** to **bound** operation

```
root@kali:~# cat /var/log/syslog | grep -i -A1 preinit
Jan  9 22:40:05 kali NetworkManager[2224]: <info> (eth0): DHCPv4 state changed
Jan  9 22:40:05 kali NetworkManager[2224]: <info>   address 192.168.198.128
```

# rotating log files

- Server logfiles fill up quickly
- To keep old logs as long as possible we use **rotation**
- when a log is full or expired a **copy** is made.
- When the copy limit is reached, new logs **overwrite** old ones
- Typically busy server logs only survive 14 days.

# Logrotate

- /etc/**logrotate**.conf looks after two log files
  - wtmp for logins
  - btmp for bad (failed) logins
- other apps look after their own logs in etc/logrotate.d

```
$ls /etc/logrotate.d
apache2  apt    lxd          rsyslog  ufw
apport   dpkg   mysql-server samba
```

- apache2 is typical

```
$cat /etc/logrotate.d/apache2
/var/log/apache2/*.log {
    weekly
    missingok
    rotate 13
    compress
    delaycompress
    notifempty
```



# sample logrotate activity

```
$sudo logrotate -d /etc/logrotate.d/apache2  
reading config file /etc/logrotate.d/apache2
```

-d = debug

Handling 1 logs

rotating pattern: /var/log/apache2/\*.log weekly (13 rotations)

empty log files are not rotated, old logs are removed

considering log /var/log/apache2/access.log

log does not need rotating

considering log /var/log/apache2/error.log

log does not need rotating

considering log /var/log/apache2/other\_vhosts\_access.log

log does not need rotating

not running prerotate script, since no logs will be rotated

not running postrotate script, since no logs were rotated



# zipped log files

- often a log file is compressed to save space
- access.log.2.gz
- need to unzip the file using **gunzip** before using cat.

```
ubuntu$ls
access.log      access.log.3.gz
access.log.1    access.log.2.gz
ubuntu$sudo gunzip access.log.2.gz
ubuntu$ls
access.log      access.log.3.gz
access.log.1    access.log.2
```

# tarballs

- A **tarball** is a group or archive of files that are bundled together using the tar command and have the .tar file extension.
- If your tar file is **compressed** using a gzip compressor, use this command to uncompress it.

```
$ tar xvzf file.tar.gz
```

Where,

x: This option tells tar to extract the files.

v: The “v” stands for “verbose.” This option will list all of the files one by one in the archive.

z: The z option is very important and tells the tar command to uncompress the file (gzip).

f: This options tells tar that you are going to give it a file name to work with.

# Shell logs

- An intruder or suspect may login and open a bash shell.
- Here she may run shell scripts with malicious intent
- In this case her activities are recorded in her `.bash_history` file
- Remember shell history in the basic cmds Lab?
  - Part 5 History and Exit

# bash shell history log

```
The suspect logs in
The suspect types the following
----
ls
date
./getinfo.sh
cat info.txt
whoami
----
The suspect logs out
----
We look at the suspect's shell history file
cd /home/group11
cat .bash_history
We see all
```

# FIN

- Bring an empty USB for the Week 9 Lab
- arrivederci